Early embryo differentiation in zebrafish

Samuel Pastva

October 26, 2026

1 Glossary and preliminaries

Because biology likes to use big, fancy words, we first have to build a glossary of terms that will appear in this report. Most are relatively "standard" in the literature, but many have slightly different meaning across subfields, so this section can be also viewed as term definitions.

- A **transcription factor (TF)** is a protein that directly influences the transcription process, either by participating in it, or by modifying the environment that influences the actual transcription machinery (e.g., opening or closing chromatin). Note that this is intentionally broader than the strict textbook definition, which only covers proteins that bind specific DNA sequences; here, we also include co-factors and chromatin remodelers.
- Time points in zebrafish embryo development are usually denoted in **hours post fertilization (hpf)**. We will shorten this to **hours (h)** when clear from context. We will be mostly interested in the development processes taking place until approx. **12hpf**.
- Since zebrafish are simple organisms, the embryo development is also relatively stable and easy to reproduce. The most important environmental factor is **temperature**, which is typically **28.5 °C**. Other temperatures within acceptable rates should have roughly linear influence on the development speed [*citation needed*]. Other factors like pH, medium density or oxygen levels also impact development, but the embryo is robust to small, reasonable variations in them [*citation needed*].
- An embryo cell can be **totipotent**, **pluripotent**, or **multipotent**. Totipotent means it can differentiate into internal embryo cells, or external cell layers enveloping the embryo (**extra-embryonic cells**). Pluripotent means it can differentiate into any embryo cell, but not into an extra-embryonic cell. Multipotent means the cell can differentiate within a specific lineage (e.g., endoderm), but not outside of that specific lineage (at least not without a significant intervention).
- We can also distinguish cells by how *robustly* they are differentiated so far. Note that these terms are originally defined based on what happens to the cell across various environments, not necessarily by the epigenetics and transcriptome, so we are "extrapolating" a bit here to what is causing these terms to be well defined in the first place.
  - A cell is **primed** if its epigenetics admit a particular differentiation event (relevant chromatin is open, important TF is expressed but not causing transcription yet, etc.). A cell can be primed towards many lineages at the same time. Importantly, a primed cell will not differentiate without some additional "signal" triggering the differentiation.
  - A cell is **specified** if it is actively transcribing RNA of a particular cell type, and will continue this process in a neutral environment. In other words, the transcription has been started and is somewhat self-sustaining, but can be easily disrupted by an unexpected "signal".

  - A cell is **determined** if it is specified, but it can maintain its own identity when encountering an opposing, but reasonable "signal" (i.e., a signal that would naturally occur somewhere in the organism). Determination can be driven by a stronger self-sustaining GRN feedback loop, as well as by epigenetic changes.
  - A cell is **differentiated** if it is determined, and it has undergone morphological changes associated with its cell type, including additional epigenetic changes to further increase robustness against signaling perturbations.
  - The expression profile of a specified / determined / differentiated cell can be different, despite being "the same cell type". The cell is actively producing TFs that will lead to epigenetic changes required for determination / differentiation, but this transcription can eventually turn off once the cell is "sufficiently committed".
  - The process is hierarchical (cell type A differentiates into B, which turns into C, ...), and for the intermediate cell types, the cell does not have to pass all stages. For example, pluripotent cells will never be truly differentiated, or even determined. Similarly, stages can interleave across cell types, so a cell can be determined type A, but specified type B, and primed type C, meaning it will maintain identity A on its own, is already developing expression profile of type B, and is able to eventually become type C if a signal is given.

- A zebrafish embryo goes through several "high-level" developmental stages: **zygote (0-0.75h)**, **cleavage (0.75-2.25h)**, **blastula (2.25-5.25h)**, and **gastrula (5.25-10h)**.
  - In the zygote and cleavage stage, the embryo looks mostly uniform. However, it is not fully symmetric: some maternal factors (dorsal determinants, germ plasm) are already localized asymmetrically (see MBT below).
  - During blastula, the surface layer of extra-embryonic cells is starting to split off. These are the **enveloping layer** (EVL) and **yolk syncytial layer** (YSL); one guards the embryo from the environment, and one is an interface with the yolk on which the embryo rests. Remaining cells are also called **deep cells**.
  - A **midblastula transition (MBT)** occurs around **2.75hpf**. Until that point, the cells are predominantly driven by the maternal transcriptome present in the zygote (a small first wave of zygotic transcription already starts around the 64-cell stage). Until MBT, the whole cell population is essentially synchronous. Nevertheless, several important decisions are already made in this period, driven by maternal factors:
    * formation of the YSL and EVL (YSL forms at 512-1024 cells; Figure 1);
    * specification of a small group of **primordial germ cells**, which inherit maternal **germ plasm** during cleavage;
    * the **dorsal-ventral** axis (see below), determined by maternal dorsal determinants that are transported to one side of the embryo.

    The causes of these decisions mostly can't be observed through RNA expression (only the consequences), since they rely on localization of maternal proteins. Also, *the overall mass of the embryo does not change much* during blastula and gastrula (the cells mostly thin out and spread); actual growth comes later. After MBT, the cells divide asynchronously, hence there is no longer a clear cell doubling pattern.
  - During gastrula, the embryo splits along the *radial axis* (the deep–superficial direction, perpendicular to the embryo surface; not to be confused with the radial symmetry of the blastula around the animal-vegetal axis) into **hypoblast** (develops into **mesendoderm**) and **epiblast** (develops into **ectoderm**). Note that hypoblast and epiblast are purely positional terms: Both are part of the **germ ring**, with hypoblast being the inner layer (adjacent to the yolk), while epiblast is the outer layer (see also

Figure 1). In other words, a cell can be a part of hypoblast and of mesendoderm; one is the location, the other is the expression profile. The germ ring is the **margin** of the embryo (you can see cells called "marginal", meaning they belong to this region). The process of separation into hypoblast/epiblast along the germ ring is called **internalization**. During gastrula, the epiblast still contains mesendoderm precursors that internalize later; only once internalization finishes does epiblast fully correspond to ectoderm.

  - Within the gastrula phase, the embryo also visibly splits along the **dorsal-ventral** axis (dorsal is back / spine, ventral is front / belly). This split is marked by a **shield** formation on the dorsal side of the embryo. Note that the dorsal side itself is determined much earlier (maternally, see MBT above); the shield is where it first becomes morphologically visible. This is again mostly a positional term (i.e., shield is not an expression profile). However, cells in the shield do produce various signaling molecules that reinforce the dorsal-ventral split, and later also the anterior-posterior separation (see below). Sometimes, they are called **organizer cells**.
  - The term **lateral** (and its opposite, **medial**) is used in two related senses. During gastrula, it refers to position along the margin: viewed from the animal pole, the margin is a circle with the dorsal side (shield) on one end and the ventral side on the opposite end; the *lateral* margin are the two sides in between (left and right). In the body, it refers to the distance from the midline (the future spine): *medial* is close to the midline, *lateral* is far from it. The two are connected, because during gastrula the cells converge towards the dorsal side, which becomes the midline of the body. A related term is **paraxial**, meaning directly alongside the midline (i.e., the body axis).
  - After gastrula, the embryo undergoes **early segmentation**, where the split along the **anterior-posterior** (A-P) axis becomes morphologically visible (as with the other axes, A-P patterning, e.g., of the neural plate, already starts during gastrula). The A-P axis is often confused with the **animal-vegetal** axis, but these are different axes: animal-vegetal is the axis of the egg (animal pole = top, where the blastoderm sits; vegetal pole = bottom of the yolk), while A-P is the body axis, which only forms during gastrula. They are only roughly aligned (the animal region gives rise to anterior head structures, while the tail bud forms near the vegetal pole). In Figure 1, AP/VP mark the animal/vegetal pole, while Ant./Pos. mark anterior/posterior. Segmentation is beyond the scope of this document.

- During blastula and gastrula, the embryo also undergoes **epiboly**, during which it envelops the connected yolk. This is again largely used as a timing / position marker, with Figure 1 showing the individual stages. Note that the formation of a hypoblast and epiblast, as well as of a germ ring and shield are happening relatively quickly between 50% and 60% epiboly, but this is a morphological change; transcriptomic and epigenetic changes leading up to the formation are already happening sooner.

- Another note is that in terms of morphology and transcriptome, epiblast is essentially just the continuation of the existing deep cells, without major changes at the moment of separation. It is the hypoblast that represents a new lineage that is splitting off of epiblast in the marginal region. However, later during gastrula, the ectoderm itself splits into **neural** and **non-neural** ectoderm (the latter becoming the **epidermis**), driven by a signaling gradient along the dorsal-ventral axis, and this split is visible in the transcriptome.

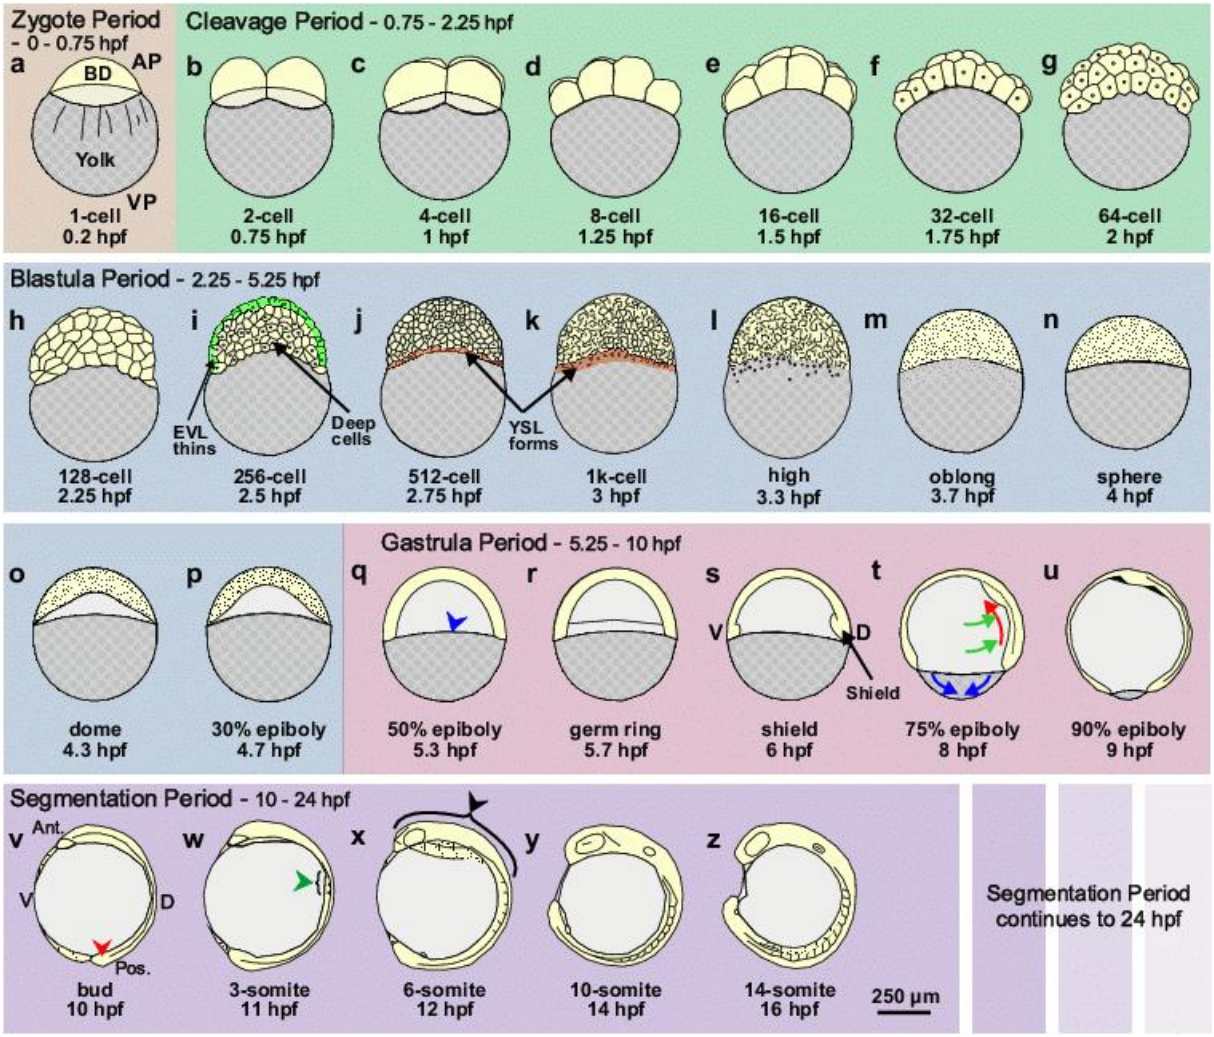


Figure 1: `https://www.aups.org.au/Proceedings/38/43-51/`

2 Transcriptomic and epigenetic timeline

In the previous section, we described the important terms and formations which appear during embryo development. However, these are mostly *morphological* terms dictating what the embryo looks like or what is happening externally. They do not capture how the cells are evolving internally, which typically precedes morphological changes quite a bit. Also, as we discussed, these do not map directly to any meaningful definition of "cell types" (i.e., a "shield cell" is not a well defined cell type). In this section, we thus summarize the *expression programs* that we expect to observe in the transcriptome during the first 12 hours, roughly when they emerge or diminish, what they do, and which earlier program they develop from. A few general remarks:

- All time ranges are approximate. They denote when a program should become *distinguishable in the transcriptome*, which usually precedes the corresponding morphological change by an hour or more. Also, the transitions are gradual: a new program typically first appears as a weak, graded bias within its parent population, and only later becomes a clearly separated cluster.
- Programs rarely "end"; most of them keep splitting into more specific sub-programs. When we say a program diminishes, we mean it is no longer a meaningful description of the cells, because they have been replaced by its more specific descendants.
- Not everything is equally visible in single-cell data. Some programs are not captured at all (e.g., the YSL, which is a syncytium, not a set of cells), or involve only a handful of cells (e.g., germ cells), and the GSE106587 atlas only starts at 3.3hpf, i.e., after MBT.

Table 1 gives an overview, and Figure 2 shows the same programs as a timeline, highlighting the branching points and the spatial axis along which each split happens. The individual programs are described in more detail below.

2.1 Cleavage and blastula (0-5.25hpf)

- **Maternal program** (0hpf, cleared by approx. 4-6hpf). Transcripts deposited in the egg by the mother. They drive the cleavage divisions and activate the zygotic genome. After MBT, a large fraction of maternal transcripts is actively degraded; by the start of gastrula, the transcriptome is dominated by zygotic expression [*citation needed*]. In the atlas, the earliest stages (high, oblong) should show a strong maternal component shared by all cells.
- **Blastomeres** (approx. 2-3hpf, diminishing 5-8hpf). The "default" program of deep cells after zygotic genome activation (ZGA). Cells are pluripotent and largely uniform; the observable differences between them are initially small and graded (distance from the margin, i.e., position along the animal-vegetal axis, and position along the dorsal-ventral axis). This program gradually disappears as the cells commit to germ layers; it persists longest in the epiblast, which is its direct continuation.
- **Yolk syncytial layer (YSL)** (2.75-3hpf, persists). Forms when marginal blastomeres collapse into the yolk cell. It drives epiboly and, importantly, is a source of signals that induce the marginal mesendoderm and contribute to the dorsal organizer [*citation needed*]. Because it is a syncytium inside the yolk, *it is not captured in single-cell data*; we only see its effect on neighbouring cells.
- **Enveloping layer (EVL)** (approx. 3-4hpf, persists as **periderm**). Develops from the outermost blastomeres, which become lineage-restricted during the blastula [*citation needed*]. Forms a protective epithelium around the embryo. In single-cell data, it is expected to be the first clearly separate cluster, already at the earliest atlas stages.

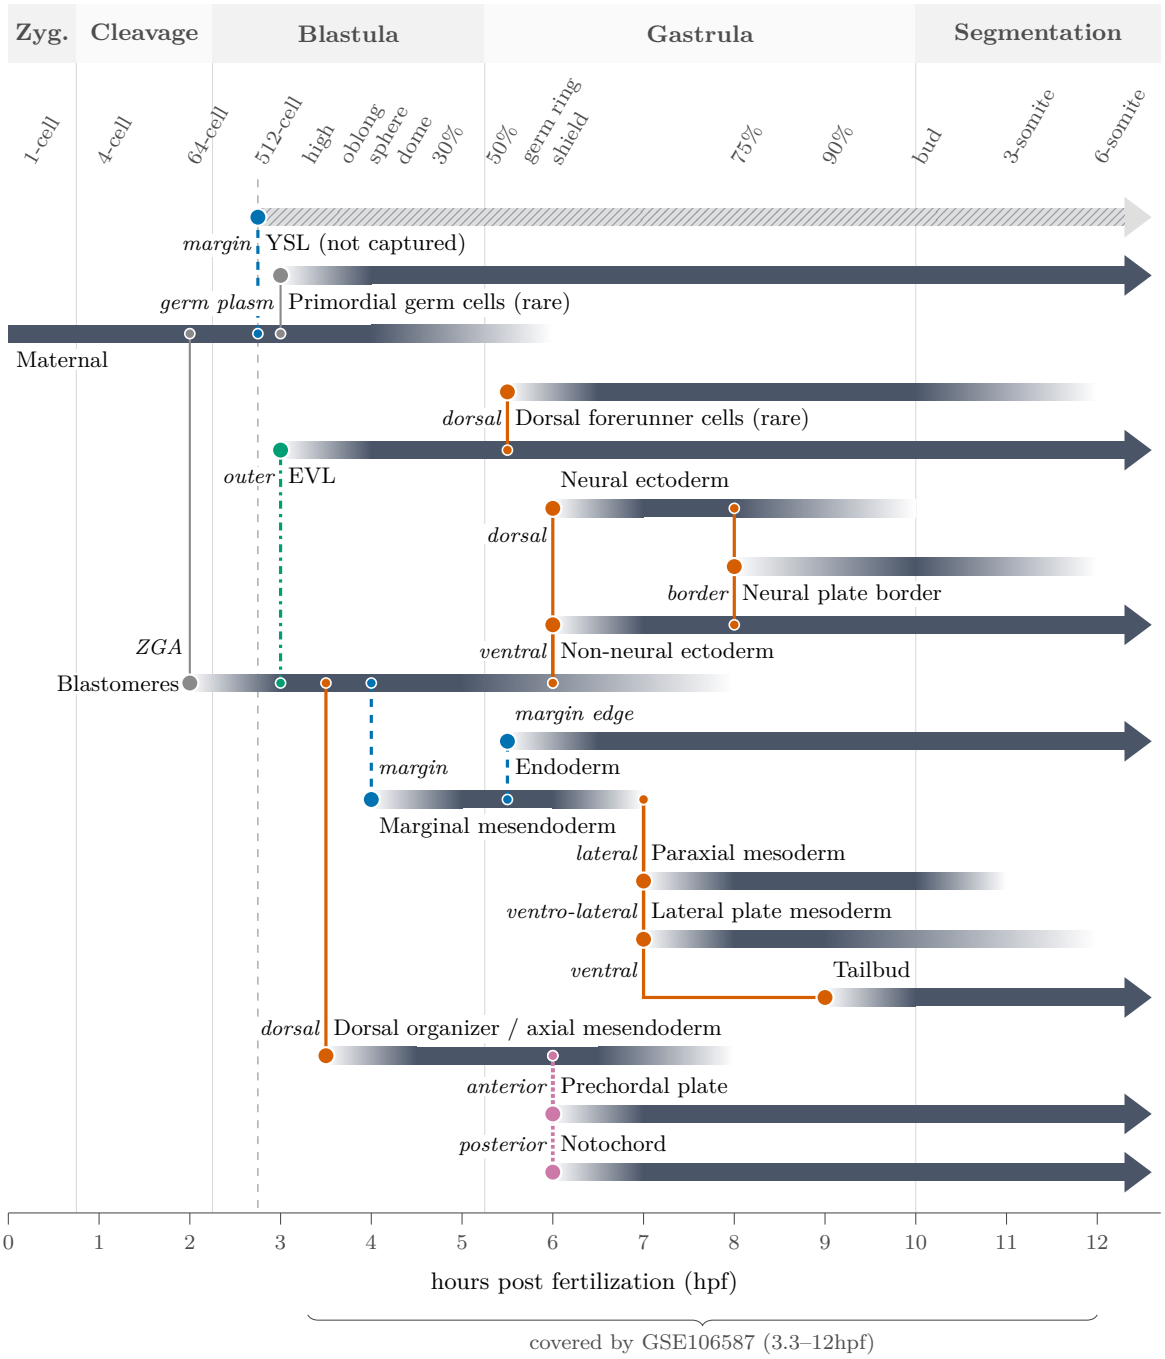


Figure 2: Timeline of the expected expression programs (Table 1). Each bar is one program; shaded ends mark the (approximate) intervals in which it emerges or diminishes. The dashed vertical line marks the MBT.

Program	Emerges (hpf)	Diminishes (hpf)	Develops from	Conditions	Note
Maternal	0	4-6	—	deposited in the egg	actively degraded after MBT
Blastomeres	2-3	5-8	maternal	zygotic genome activation (MBT)	replaced by germ layer programs; persists longest in the epiblast
YSL	2.75-3	—	maternal	marginal blastomeres fusing with the yolk	syncytium; not captured in single-cell data
EVL	3-4	—	blastomeres	outermost cell layer	becomes periderm
Primordial germ cells	3-4	—	maternal	inherited germ plasm	rare (a handful of cells)
Dorsal organizer	3.5-4.5	6.5-8	blastomeres	dorsal margin; maternal dorsal determinants	morphologically the shield (6hpf); splits into prechordal plate and notochord
Marginal mesendoderm	4-5	6-7	blastomeres	margin; signaling from YSL	forms the hypoblast; splits into endoderm and mesoderm programs
Endoderm	5.5-6.5	—	marginal mesendoderm	closest to the margin; strongest signaling	few cells at this stage; gut, liver, pancreas, pharynx
Dorsal forerunner cells	5.5-6.5	10-12	EVL	dorsal margin, ahead of the shield	rare; form the Kupffer's vesicle
Prechordal plate	6-7	—	dorsal organizer	first internalized cells	anterior part becomes the hatching gland (10-12hpf)
Notochord	6-7	—	dorsal organizer	later internalized cells	
Paraxial mesoderm	7-8	10	marginal mesendoderm	lateral margin; further from the margin edge	splits into adaxial cells and presomitic mesoderm; segments into somites
Lateral plate mesoderm	7-8	9-12	marginal mesendoderm	ventral and lateral margin	splits into heart, blood, endothelial and kidney precursors
Non-neural ectoderm	6-7	—	blastomeres	ventral epiblast; ventralizing signaling	becomes epidermis
Neural ectoderm	6-7	8-10	blastomeres	dorsal epiblast; protected by the dorsal organizer	subdivides along the A-P axis (forebrain, midbrain, hindbrain, spinal cord)
Neural plate border	8-10	10-12	neural and non-neural ectoderm	boundary between the two; intermediate signaling	splits into neural crest and placodes
Tailbud	9-10	—	marginal mesendoderm	ventral margin converging at the vegetal pole	includes overlying epiblast; neuromesodermal progenitors

Table 1: Overview of expected expression programs. Times are approximate and refer to when the program becomes distinguishable in the transcriptome; — means the program persists beyond 12hpf. *Develops from* is always a preceding program in the table; *conditions* describe where (location) or how (signaling, inheritance) the new program arises within it.

- **Primordial germ cells (PGCs)** (specified during cleavage, own zygotic program from approx. 3-4hpf, persists). Blastomeres that inherited maternal germ plasm. They give rise to gametes and do not participate in somatic differentiation. Only a handful of cells per embryo, so they will form a very small cluster.

- **Dorsal organizer** (approx. 3.5-4.5hpf, splits at 6.5-8hpf). Dorsal marginal cells (together with the dorsal YSL) that respond to the maternally determined dorsal side of the embryo. They produce signals that dorsalize the surrounding tissue, protect the dorsal ectoderm from ventralizing signals (enabling neural fate), and later contribute to A-P patterning. Morphologically, this becomes the **shield** at 6hpf; transcriptionally, the program is present well before that. Its cells later give rise to the axial mesoderm (prechordal plate and notochord). From approx. 6hpf, the internalized organizer cells (the inner layer of the shield) are often called the **axial mesendoderm**, i.e., midline mesendoderm that has not yet split into prechordal plate and notochord. We do not treat it as a separate program, because it is transcriptionally a gradual continuation of the organizer program; the clear split only happens with the prechordal plate and notochord.

- **Marginal mesendoderm** (approx. 4-5hpf, splits at 6-7hpf). Cells within a few cell tiers of the margin, induced by signals from the YSL and neighbouring marginal cells. The program is graded: the closer to the margin (and the stronger the signal), the more the cells are biased towards endoderm, while cells further away are biased towards mesoderm [*citation needed*]. These are the cells that internalize to form the hypoblast.

2.2 Gastrula (5.25-10hpf)

- **Endoderm** (approx. 5.5-6.5hpf, persists). Develops from the marginal mesendoderm cells that received the strongest inducing signal (mostly close to the margin, more on the dorsal side). Will form the gut and associated organs (liver, pancreas) and the pharynx. Relatively few cells at this stage.

- **Dorsal forerunner cells** (approx. 5.5-6.5hpf, forming the **Kupffer's vesicle** at approx. 10-12hpf). A small group of cells at the dorsal margin, ahead of the shield, partially derived from the EVL [*citation needed*]. They do not internalize; instead, they migrate ahead of the margin and form the Kupffer's vesicle, which establishes the left-right asymmetry. Only a handful of cells.

- **Prechordal plate** (approx. 6-7hpf, persists). The anterior part of the axial mesoderm, derived from the leading (first internalized) cells of the dorsal organizer. It migrates towards the animal pole and contributes to head patterning. Its most anterior part becomes the **hatching gland** (approx. 10-12hpf).

- **Notochord** (approx. 6-7hpf, persists). The posterior part of the axial mesoderm, derived from the remaining dorsal organizer cells. Forms the midline rod that later supports the body axis, and is an important source of signals that pattern the neighbouring neural tube and paraxial mesoderm.

- **Paraxial mesoderm** (approx. 7-8hpf, starts to form somites at 10hpf). Mesoderm from the lateral margin that converges on both sides of the notochord. Splits into **adaxial cells** (directly adjacent to the notochord; slow muscle precursors) and **presomitic mesoderm**, which will be segmented into somites (muscle, skeleton).

- **Lateral plate mesoderm** (approx. 7-8hpf, splitting at 9-12hpf). Mesoderm from the ventral and lateral margin. Gives rise to the heart, blood and endothelial precursors, and the pronephros (kidney). Many of these sub-programs only become clearly distinct around the end of gastrula or early segmentation.

- **Non-neural ectoderm** (approx. 6-7hpf, persists as **epidermis**). The ventral part of the epiblast, where ventralizing signals are strong. It is the first ectodermal program to diverge from the blastomere program.
- **Neural ectoderm** (approx. 6-7hpf, subdivided at 8-10hpf). The dorsal part of the epiblast, protected from ventralizing signals by the dorsal organizer. Forms the **neural plate**. During late gastrula, it is patterned along the A-P axis into forebrain (including eye field), midbrain, hindbrain and spinal cord programs.
- **Neural plate border** (approx. 8-10hpf, splitting at 10-12hpf). Cells at the boundary between neural and non-neural ectoderm, receiving intermediate levels of signals. Gives rise to the **neural crest** (very diverse lineage: pigment cells, craniofacial skeleton, peripheral neurons, ...) and **placodes** (sensory organs).
- **Tailbud** (approx. 9-10hpf, persists). Forms from ventral marginal cells converging at the vegetal pole at the end of epiboly. Contains bipotent **neuromesodermal progenitors** that produce both spinal cord and paraxial mesoderm during body axis elongation. Note that this is one of the few programs which mix "neural" and "mesodermal" signatures.

3 Dataset overview

The analysis is based on the GEO SuperSeries GSE106587 [2]. It consists of three single-cell datasets:

- **Drop-seq atlas** (SubSeries GSE106474). Wild-type embryos at 12 stages from high (3.3hpf) to 6-somite (12hpf); 39,505 cells, 23,974 genes. This is the main dataset.
- **SMART-seq2** (SubSeries GSE106466). 416 cells at 50% epiboly: 52 wild-type cells and 364 cells of MZ*oep* mutants, which lack Nodal signalling and hence most of the mesendoderm. Fewer cells, but deeper coverage per cell.
- **10x Genomics** (GSE106587 itself). Wild-type (3,000 barcodes) and MZ*oep* (2,200 barcodes) embryos at 6-somite (12hpf). The barcode lists are not QC-filtered, so not every barcode is a cell.

All three are available only as processed count matrices (raw reads were not downloaded). In the rest of this section, we only describe the Drop-seq atlas. **The GEO matrix has already passed the basic quality filters of the original study, but not its later outlier removal** (Section 3.4); here, we use all 39,505 cells without any additional filtering. The experimental details below are taken from the supplementary materials of the paper. The numbers, figures and table are produced by `scripts/01_dataset_structure.py`.

3.1 Stages, batches and libraries

The Drop-seq atlas has three nested levels (Figure 3 and Table 2):

- **Stage**: the developmental time point (12 stages, each with a single hpf value).
- **Sequencing batch** (DS2, DS3, DS4 and DS5): the "sequencing batch" field of the GEO samples. It is the main expected source of technical variation.
- **Library**: one sample of pooled embryos, processed in one Drop-seq run; equivalently, one GEO sample (GSM) and one SRA experiment. There are 28 libraries in total. Most (stage, batch) pairs have exactly one library. Six pairs have two, in which case the cell names of the second library carry a `b` suffix (e.g., `ZF50_WT_DS4b_...`).

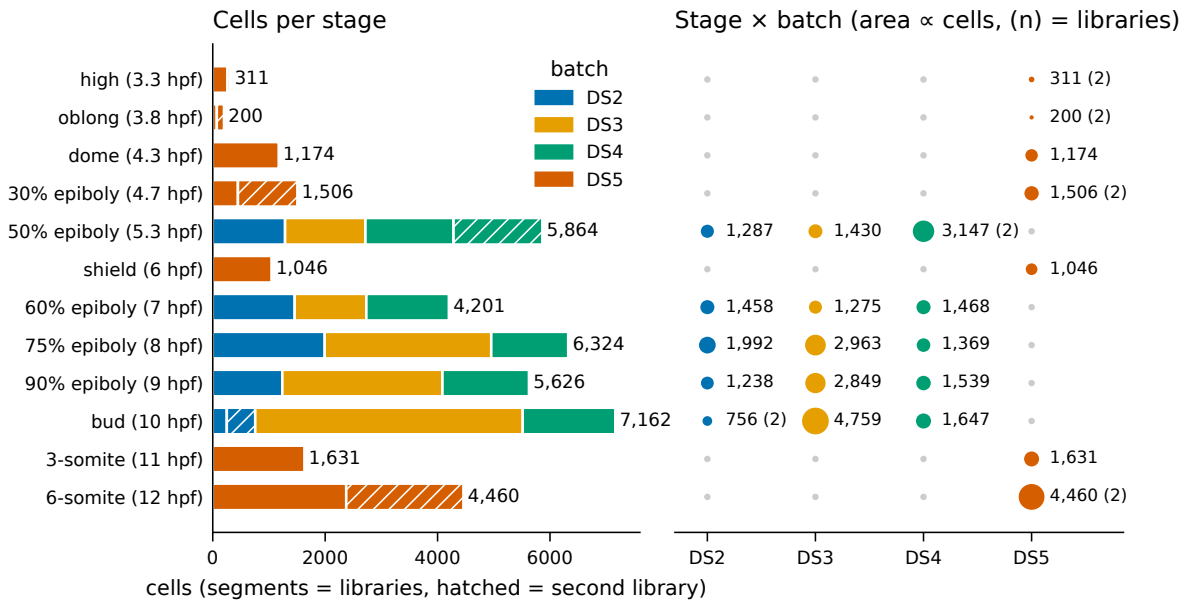


Figure 3: Structure of the Drop-seq atlas (GSE106474). Left: number of cells per stage; each segment is one library, coloured by sequencing batch, and second libraries of the same stage and batch are hatched. Right: the same counts per stage and batch. Stages of batch DS5 and stages of batches DS2–DS4 do not overlap.

Each library was prepared from a pool of 20–40 embryos of the same stage. The embryos were manually deyolked (**which also removes the YSL**), mechanically dissociated, and the resulting cell suspension was encapsulated in Drop-seq droplets. Consequently, cells carry no information about which embryo (or which position in the embryo) they came from. Some libraries were sequenced twice to increase depth, but the reads of such re-sequencing are merged into the same library; they do not form a second library.

Second libraries are separate embryo pools. GEO itself does not say whether the second library of a (stage, batch) pair comes from a separate embryo pool (biological replicate) or is a second run of the same pool (technical replicate). However, the supplementary Table S1 of the paper lists the total number of embryos used for each stage (Table 3). For high (56 embryos), oblong (76) and 30% epiboly (59), this is more than a single pool of at most 40 embryos, so the two libraries of these stages must come from different pools. For the remaining stages with a second library (50% epiboly, bud and 6-somite), the embryo counts are consistent with separate pools of approx. 20 embryos each, but they do not rule out a shared pool. The paper also treats each library as its own unit: Table S1 calls the libraries "batches" (two for high, four for 50% epiboly), and the two 6-somite libraries are described as "acquired in two batches" and batch-corrected before trajectory reconstruction. We thus treat second libraries as biological replicates sequenced in the same sequencing batch. Note that the paper's "batch" corresponds to our library, not to the GEO sequencing batch.

GEO replicate numbering. The "Rep N" in the GEO sample titles numbers the libraries of a stage *across batches*. It does not correspond to our a/b label. For example, "50% epiboly Rep 2" is the DS3 library, while "high Rep 2" is the second DS5 library of the high stage.

Table 2: Drop-seq libraries. Lib. "b" is the second library of the same stage and batch (a separate embryo pool). Medians are per cell, over all cells in the GEO matrix.

Stage	hpf	Batch	Lib.	GEO sample	Cells	Median UMIs	Median genes
high	3.3	DS5	a	GSM2838529	263	6,731	3,009
		DS5	b	GSM2838530	48	13,385	4,566
oblong	3.8	DS5	a	GSM2838531	72	11,866	4,171
		DS5	b	GSM2838532	128	3,603	2,048
dome	4.3	DS5	a	GSM2838528	1,174	4,398	1,931
30% epiboly	4.7	DS5	a	GSM2838506	448	3,858	1,717
		DS5	b	GSM2838507	1,058	3,260	1,564
50% epiboly	5.3	DS2	a	GSM2838509	1,287	5,600	2,106
		DS3	a	GSM2838510	1,430	4,562	1,730
		DS4	a	GSM2838511	1,567	5,179	1,988
		DS4	b	GSM2838512	1,580	5,630	1,962
shield	6	DS5	a	GSM2838533	1,046	2,904	1,291
60% epiboly	7	DS2	a	GSM2838513	1,458	6,173	1,898
		DS3	a	GSM2838514	1,275	6,074	1,826
		DS4	a	GSM2838515	1,468	4,970	1,640
75% epiboly	8	DS2	a	GSM2838518	1,992	5,404	1,768
		DS3	a	GSM2838519	2,963	3,352	1,258
		DS4	a	GSM2838520	1,369	5,582	1,712
90% epiboly	9	DS2	a	GSM2838521	1,238	5,814	1,793
		DS3	a	GSM2838522	2,849	3,377	1,282
		DS4	a	GSM2838523	1,539	3,603	1,343
bud	10	DS2	a	GSM2838524	250	6,642	1,915
		DS2	b	GSM2838525	506	6,446	1,864
		DS3	a	GSM2838526	4,759	2,556	1,063
		DS4	a	GSM2838527	1,647	3,683	1,381
3-somite	11	DS5	a	GSM2838508	1,631	2,609	1,047
6-somite	12	DS5	a	GSM2838516	2,375	2,000	884
		DS5	b	GSM2838517	2,085	3,180	1,237

3.2 Batch design

The most important property of the atlas is that the batches split the stages into two blocks that *share no stage*:

- **Block A** (batch DS5 only): high, oblong, dome, 30% epiboly, shield, 3-somite and 6-somite; 11 libraries, 10,328 cells.
- **Block B** (batches DS2, DS3 and DS4): 50% epiboly, 60%, 75% and 90% epiboly, and bud; 17 libraries, 29,177 cells. Every stage of this block was sequenced in all three batches.

Within block B, the design is balanced, so stage and batch effects can be separated. Within block A, there is only one batch, so stage differences are not confounded with batch, but there is also nothing to estimate the DS5 batch effect from. This is particularly unfortunate for our focus on the decisions around 6hpf. Shield (6hpf, DS5) lies between 50% epiboly (5.3hpf) and 60% epiboly (7hpf), which are both in block B. A batch effect will therefore appear as an abrupt developmental change at shield and back again at 60% epiboly.

3.3 Library size and sequencing depth

The number of cells per stage ranges from 200 (oblong) to 7,162 (bud) (Figure 3). The five block B stages (50% epiboly to bud) contain 74% of all cells. The two earliest stages are tiny: high and oblong have 511 cells together, and three of their four libraries have fewer than 150 cells (48, 72 and 128). Sequencing depth (Figure 4) is a median of 3,798 UMIs and 1,444 detected genes per cell over the whole atlas. We can make the following observations:

- **Depth decreases with developmental time**, from library medians of 4–13 thousand UMIs per cell at high and oblong to 2–3 thousand at 3- and 6-somite. Part of this is biological (cells get smaller and contain less mRNA with each division, which is also why the original study used stage-specific filtering thresholds).
- **Libraries of the same stage and batch are not interchangeable.** At oblong, the median depth of the first library is 3.3× that of the second; at high, the second library is 2× deeper than the first, and at 6-somite 1.6×.

3.4 Cell filtering

The original study filtered the cells in two steps; the dataset includes the first, but not the second filtering.

Basic quality filters (already applied in GEO). First, cells were filtered by the number of detected genes and UMIs, using a separate range for each stage (Table 3), because RNA content per cell decreases over early development. Cells below the range are likely low-quality (e.g., broken cells or empty droplets with ambient RNA), while cells above the range are likely doublets. Afterwards, cells with more than 45% mitochondrial RNA (a sign of stressed or dying cells) were removed. The GEO matrix is consistent with these filters having been applied:

- Only 51 of the 39,505 cells fall outside the gene ranges, all of them just slightly.
- The mean number of detected genes per stage agrees with Table S1 to within 1.5%.
- The GEO matrix contains no mitochondrial genes. This explains why the mean UMIs per cell are 20–25% lower than in Table S1 (where approx. 20–29% of UMIs are mitochondrial), and why some cells fall below the UMI ranges.

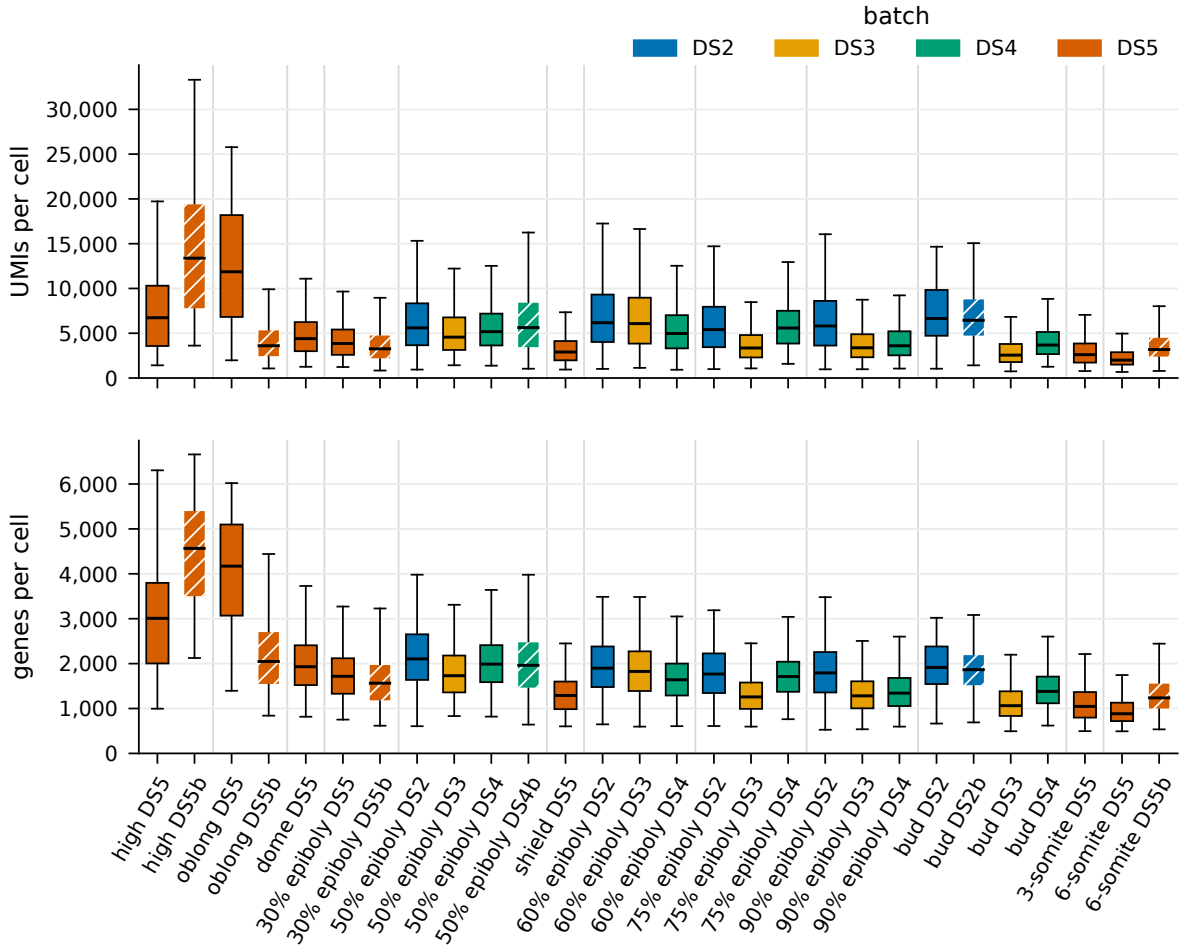


Figure 4: Per-library sequencing depth: UMIs (top) and detected genes (bottom) per cell. Boxes show the quartiles, whiskers extend to 1.5× the interquartile range (outliers not shown). Colour is the sequencing batch; second libraries are hatched. Thin vertical lines separate stages.

Outlier removal (not applied in GEO). Before the trajectory reconstruction, the paper removed 618 cells forming a cluster dominated by an "apoptotic-like" expression program, which appears in cells of many different types and would therefore create artificial connections between unrelated cell types. It also removed 139 cells that were poorly connected in the nearest-neighbour graph. The paper's full-data visualization uses 39,488 cells, and $39{,}488 - 618 - 139 = 38{,}731$, which is the number of cells reported in the paper. The GEO matrix contains 17 cells more than 39,488, which the paper does not explain. The identities of the removed cells are not available in GEO. Finally, the per-stage cell counts in Table S1 sum to 38,671, i.e., 60 fewer than 38,731, so the paper is not fully consistent here either.

3.5 Consequences for the analysis

- **Quantitative analysis within library, qualitative comparison across libraries.** To avoid batch and stage effects, quantitative analysis (normalization, binarization, etc.) should be calibrated and applied within individual libraries. Between libraries, we should primarily draw conclusions from qualitative results (e.g., gene A is higher than B in cells X across DS2/DS3/DS4; the quantity can be different, but the effect stays the same).

- **When to compare quantities across libraries?** We will most likely need to break the

Table 3: Cell filtering per stage. Gene and UMI ranges are the per-stage thresholds of the original study (UMIs including mitochondrial RNA). Cells (GEO) is the number of cells in the GEO matrix; cells and embryos (Table S1) are taken from the supplementary Table S1.

Stage	Genes	UMIs	Cells (GEO)	Cells (Table S1)	Embryos (Table S1)
high	1,000–7,500	1,500–40,000	311	311	56
oblong	625–7,500	1,500–30,000	200	199	76
dome	800–3,800	2,000–20,000	1,174	1,147	28
30% epiboly	625–3,000	1,000–17,500	1,506	1,454	59
50% epiboly	600–4,000	1,500–25,000	5,864	5,695	91
shield	600–2,500	1,000–15,000	1,046	1,020	22
60% epiboly	600–3,500	1,500–22,500	4,201	4,099	79
75% epiboly	600–3,200	1,400–20,000	6,324	6,170	73
90% epiboly	500–3,500	1,000–20,000	5,626	5,474	66
bud	500–3,200	1,000–17,500	7,162	7,112	85
3-somite	500–3,000	1,000–12,500	1,631	1,602	19
6-somite	500–3,000	1,000–12,500	4,460	4,388	40
total			39,505	38,671	694

first rule in a few places. First, comparing **time stages within batch** is reasonable when assessing long-term changes in transcription (but should probably carry greater burden of proof than within-library tests). Similarly, **batch-effect calibration within stage** is possible for DS2, DS3 and DS4. This is not possible with respect to DS5, but maybe we can at least extrapolate some useful information from it. Batch correction ideally stays within each stage to avoid "correcting for" biological development. But most algorithms can take this into account, so in theory, for DS2/DS3/DS4, a full batch correction run is viable; however, it is questionable whether it is needed to derive useful qualitative comparisons. Finally, **library-effect within batch and stage** can be also quantified using the a/b replicates, but we only have these for a small number of libraries, and they do come from separate embryo pools, so there can be genuine biological differences between them.

- **Some pseudo-libraries may be needed.** The early development stages (high, oblong) contain a very small number of cells, further subdivided into the a/b replicates. Assuming we don't observe significant expression profile differences between the a/b replicates, it may be reasonable to merge them and treat them as one library to obtain a bigger cell pool for the following experiments. In this context, **down-sampling or weighting the remaining, much larger libraries may be necessary** for some analysis methods to avoid over-representation of the 50%/60%/75%/90%/bud stages.

- **Outliers must be handled by us.** The GEO matrix still contains the apoptotic-like cells and the poorly connected outliers removed in the paper. We likely need our own (comparable) outlier filtering step. Furthermore, some cell types not relevant for the Boolean model may be also eliminated outright (e.g. the germ cells, or the EVL). However, for operations such as discretization, it may be useful to have the whole expression distribution, even if it is not fully assigned to cell types.

4 Result: a/b libraries have distinct expression profiles

Six (stage, batch) pairs have two libraries, a and b (Table 2). If the two libraries of a pair were interchangeable, we could merge them, which would remove one level of the library hierarchy and, in particular, give larger cell pools for the tiny high and oblong stages. In this section, we test whether the two libraries of a pair contain the same population of expression profiles.

They do not: in every pair, the libraries can be told apart, often as clearly as two neighbouring stages, and neither depth normalization nor depth matching changes this. We therefore keep the a/b libraries separate. The numbers and figures are produced by `scripts/02_replicate_mixing.py` (mixing) and `scripts/03_replicate_differences.py` (gene-level evidence).

4.1 Experiment

For each pair, we embed the cells of both libraries together using a standard pipeline: each cell is normalized to 10,000 UMIs and log-transformed, the 2,000 most variable genes are selected on both libraries jointly, and the cells are projected onto the first 30 principal components using PCA. In this space, we find the 20 nearest neighbours of every cell. If the libraries contain the same population of cells, the neighbourhood of a cell is a random sample of both libraries, i.e., the number of neighbours from library b follows a binomial distribution around the overall share of b cells. If one library contains cells that the other lacks, or if all cells of one library are shifted in expression, the neighbourhoods are dominated by one library. We summarize this with three measures; the first two are standard metrics of batch mixing in single-cell data [6]:

- **Rejection rate** (in the style of kBET [1]): the fraction of cells whose neighbourhood composition differs from the overall composition (binomial test, $p < 0.05$).
- **Local inverse Simpson's index** (LISI [4]): the effective number of libraries in a neighbourhood (between 1 and 2). It is computed as $1/(p_a^2 + p_b^2)$, where p_a and p_b are the fractions of neighbours from each library, so a 50/50 neighbourhood gives 2 and a single-library neighbourhood gives 1. For simplicity, we weight the 20 neighbours uniformly instead of using the original distance-based weights.
- **Classifier AUC**: a logistic regression predicts the library of a cell from its principal components (5-fold cross-validation). This is a classifier two-sample test [5], a general statistical method rather than an established single-cell metric. A related idea was used by the authors of the atlas, who identified batch-specific gene modules by how well they predict the batch (ROC AUC, via Seurat's `BatchGene` function) [2]. An AUC of 0.5 means the libraries are indistinguishable, 1.0 means they are fully separable.

None of these has a natural threshold for "mixed well", and with thousands of cells, even a negligible difference becomes statistically significant. We therefore calibrate every pair against reference comparisons of the same size:

- **Permutation null**: the a/b labels are randomly shuffled (200 distinct repetitions) on the same neighbourhood graph. This is what perfect mixing looks like.
- **Negative control**: a random split of one library into two parts of the same sizes as the pair, analysed with the whole pipeline. This is what two samples of the same embryo pool look like.
- **Positive control**: library a against an adjacent stage of the same batch (e.g., 50% epiboly DS4 against 60% epiboly DS4). This is what a real developmental difference looks like.
- **Other batch** (where available): library a against the same stage from another batch (50% epiboly DS3, bud DS3).

Finally, we know that the libraries of a pair differ in sequencing depth (up to 3.3× in the median, Figure 4). Normalizing each cell to the same total removes differences in the average expression level, but not in sparsity: a shallow cell has more zeros, which can by itself move it in the embedding. We thus run every comparison twice:

- **Normalized**: all UMIs, only the per-cell normalization described above.
- **Thinned**: before normalization, each UMI of the deeper library is kept with probability equal to the ratio of the two median depths (binomial thinning [3]). Afterwards, both libraries have the same median depth, as if the deeper one had been sequenced less.

4.2 Results

The results are summarized in Figure 5. The full per-cell distributions of neighbourhood composition for all comparisons are in `results/figures/replicate_mixing_hist.pdf`, the UMAPs of all six pairs in `results/figures/replicate_mixing_umap_normalized.pdf` and `..._thinned.pdf`, and all numbers in `results/tables/replicate_mixing.tsv`.

- **The pipeline itself does not separate the libraries.** The random splits are indistinguishable from the permutation null (AUC 0.46–0.55, rejection rate at most 4%).
- **All six pairs are clearly distinguishable.** The classifier AUC is 0.83–1.00, while the highest AUC in the permutation null is 0.51–0.63. Between 14% and 81% of cells have a neighbourhood dominated by one library, compared to at most 4% under the null.
- **The difference is often comparable to a stage step.** For high and 6-somite, a against b is about as separable as a against the adjacent stage. At 50% epiboly, DS4a is as different from DS4b as from DS3 (81% vs. 81% rejected neighbourhoods). The bud DS2 pair is the most similar (AUC 0.83, 14% rejected), but there, the libraries of different batches (DS2 against DS3) are also unusually similar.
- **Depth is not the cause.** In the normalized experiment, depth alone predicts the library with AUC up to 0.76 (high). After thinning, this drops to approx. 0.50 in all pairs, but the libraries remain exactly as separable (AUC 0.85–1.00).
- **The difference is spread over many genes.** Removing the 1,000 genes with the largest a/b fold change and repeating the analysis leaves five of the six pairs separable (AUC 0.67–0.92). Only for bud DS2 does this make the libraries indistinguishable (AUC 0.52).
- **The libraries share their structure, but are shifted within it.** In the UMAPs (Figure 5, bottom), both libraries cover the same branches, but within each branch, the cells of a and b occupy neighbouring, mostly non-overlapping regions. This suggests that most of the difference is a shift of the whole library rather than a cell population unique to one of them.

Consequently, merging the a/b libraries would mix two distinct expression profiles. Normalization or thinning does not make them compatible. If we need larger cell pools for some analysis, the two libraries have to be treated like two batches (i.e., with the same care as any other between-library comparison), not as one population. This agrees with the original study, which treated libraries as batches: at 50% epiboly (four libraries), 7 of the 20 gene modules it identified reflected batch effects and were removed, and the two 6-somite libraries were batch-corrected before trajectory reconstruction [2].

4.3 Hypotheses: why do the libraries differ?

The analysis above shows *that* the libraries differ, not *why*. Here, we collect the evidence for three plausible explanations. This is not meant to be exhaustive or conclusive. It is based on a pseudobulk comparison of each pair (mean expression per 10,000 UMIs, and the fraction of cells in which a gene is detected, computed after thinning the deeper library so that detection rates

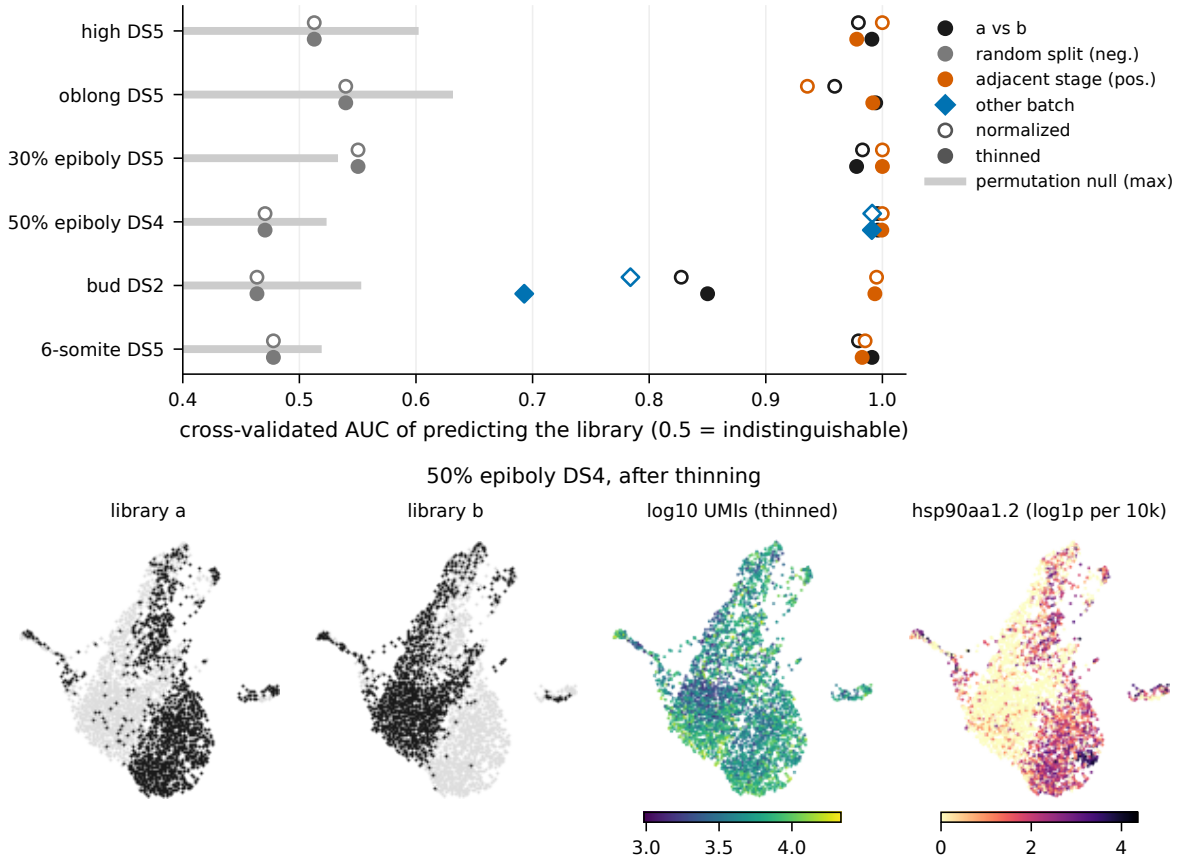


Figure 5: Mixing of the a/b libraries. Top: cross-validated AUC of a classifier predicting the library of a cell from its principal components, for each pair (a vs. b, black) and its reference comparisons; open markers without, filled markers with binomial thinning of the deeper library. The grey bar extends to the highest AUC obtained with randomly permuted labels. Bottom: UMAP of the thinned 50% epiboly DS4 pair, with separate panels for the cells of each library, the depth after thinning, and the expression of the heat shock gene *hsp90aa1.2* (Section 4.3).

are comparable). With only one library per side, there is no replicate-level variance, so we do not report p-values. Table 5 lists the genes supporting each explanation; all genes with at least a 2-fold difference are in `results/tables/replicate_de_genes.tsv`.

Different parents (clutch effect). Each library is a separate pool of embryos, quite possibly from different parents. The supplement of [2] only states that the embryos come from wild-type crosses of the Tupfel longfin/AB background, a mix of two laboratory strains, and does not say whether the pools of one stage come from different crosses or days. Wild-type zebrafish lines are genetically heterogeneous, and before and shortly after zygotic genome activation, nearly all mRNA in a cell is inherited from the mother. A genetic difference between parents should show up as *on/off genes*, detected in many cells of one library and almost none of the other, regardless of cell type. Such genes exist in high (e.g., *gbgt1l4*: 6% vs. 71% of cells) and 50% epiboly DS4 (e.g., *si:dkey-108k21.10*: 49% vs. 5%), see Table 4. At high, library b also contains approx. 2–4× more of the maternal germ plasm transcripts *nanos3* and *dnd1*, which is consistent with a different maternal contribution. Checking whether the on/off genes lie in genomic regions known for copy-number variation would strengthen this explanation.

Handling stress. In the 50% epiboly DS4 pair, library *a* strongly expresses heat shock genes (*hsp90aa1.2* 16×, *hsp70.3* 9×, *serpinh1b* 7×). *hsp90aa1.2* is detected in 81% of the cells of library *a*, but only in 14% of library *b*, and it is high throughout library *a* in the UMAP (Figure 5), i.e., this is a shift of the whole library, not a subpopulation. Heat shock genes respond within minutes to hours, so this most likely reflects the handling of the embryos or cells (temperature, dissociation) rather than development. The protocol leaves room for such differences: dissecting one sample took up to 15 minutes, cells were dissociated mechanically, and droplets were held on ice for up to 1.5 hours to allow several sequential collections before reverse transcription [2], so two libraries of the same stage and batch may have been handled for different lengths of time. Note that this also affects any comparison involving DS4*a*, e.g., heat shock genes appear among the top differences between 50% and 60% epiboly in DS4.

Developmental timing. Embryos of one stage are collected by morphology: each dish was inspected visually, and 20–40 embryos were taken once they reached the desired stage [2]. Two pools of the same nominal stage can therefore differ by some minutes of development. We use two independent indicators of developmental time, each expressed in units of one "stage step" (the difference between the pair and an adjacent stage of the same batch, so that, e.g., +0.25 means that *b* is a quarter of a stage step older than *a*):

- **Ribosomal fraction**: the share of UMIs from the 79 ribosomal protein genes. Across the atlas, it increases monotonically with developmental time, from approx. 2% at high to 21–24% at bud, and within a library, it does not correlate with the depth of a cell (median Spearman correlation −0.05; `results/tables/library_ribosomal.tsv`). However, it no longer increases consistently between 3- and 6-somite.
- **Stage-marker score**: the ratio of UMIs in genes that are more than 2× higher in the older stage to UMIs in genes that are more than 2× higher in the younger stage (in both libraries of the pair; ribosomal genes excluded). For high and oblong, the reference stage is dome, so this score measures the progress of zygotic genome activation.

We consider timing supported only if both indicators agree (Table 4). This is the case for **bud DS2, where library *a* is older** (by 0.5 and 0.2 stage steps). This is consistent with library *a* containing more cells expressing the muscle genes of the adaxial cells (*myod1*, *myl10*, *actc1a*), which start to express muscle genes during late gastrulation, and fewer cells expressing *sox19b*, an early ectodermal SoxB1 gene. The low detection rates of these genes (at most 21%) suggest that this is a difference in the number of such cells, as expected if library *a* is slightly more advanced. It is also the pair where removing the top 1,000 genes eliminates the difference, so timing may be the main explanation here. In 50% epiboly DS4 (library *b* older) and oblong (library *a* older), the indicators agree in direction, but the marker score shift is small. At high and 30% epiboly, they disagree. In particular, at high, the zygotic score suggests that library *b* is *younger*, while its ribosomal fraction is higher. At the earliest stages, development changes the ribosomal fraction by less than one percentage point, and in both the high and oblong pairs, the library with the higher ribosomal fraction is the much deeper one, so we suspect a technical component and do not trust this indicator before dome. For 6-somite, neither indicator is usable, because neither changes consistently between 3- and 6-somite.

Unexplained differences. For 30% epiboly and 6-somite, we found no specific explanation: both are separable (AUC 0.98), but none of the three explanations is clearly supported. Furthermore, this analysis cannot separate a difference in cell type composition from a shift in expression within cell types. For example, the higher expression of EVL keratins (*krt8*, *krt18*, *cldne*) in oblong *b* comes with higher detection rates (4–7% vs. 12–21%), so library *b* probably simply contains more EVL-like cells. We will revisit this distinction once the cell types are identified.

Table 4: Summary of the a/b differences. On/off genes are detected in at least 40% of the cells of one library and at most 10% of the other (after depth matching). Ribosomal % is the median share of UMIs from ribosomal protein genes. Shifts are in stage steps, positive if library b is older; n/a if the indicator does not change between the pair and the reference stage. Timing is "a/b older" if both shifts agree and are at least 0.1, "weak" if they agree but one is smaller, and "conflicting" if they disagree. Reference stages: dome (high, oblong), shield (30% epiboly), 60% epiboly DS4 (50% epiboly), 90% epiboly DS2 (bud) and 3-somite (6-somite).

Pair	Cells (a/b)	On/off	Ribo. % (a/b)	Ribo. shift	Marker shift	Timing
high DS5	263/48	2	1.8/2.7	+0.58	−0.14	conflicting
oblong DS5	72/128	2	2.9/2.0	−0.70	−0.04	weak, a older
30% epiboly DS5	448/1,058	0	6.6/6.5	−0.10	+0.07	conflicting
50% epiboly DS4	1,567/1,580	3	7.8/9.1	+0.33	+0.10	weak, b older
bud DS2	250/506	0	24.3/21.3	−0.49	−0.20	a older
6-somite DS5	2,375/2,085	0	29.4/25.9	n/a	n/a	n/a

Table 5: Selected genes supporting the explanations of the a/b differences. Mean expression is in UMIs per 10,000 UMIs; the fold change is b relative to a. Detection is the fraction of cells with at least one UMI, after thinning the deeper library to the median depth of the other.

Pair	Gene	Reading	Mean/10k (a/b)		$\log_2$FC	Detected (a/b)	
high DS5	*gbgt1l4*	on/off	0.1	2.3	+3.5	6%	71%
	si:dkey-261j4.3	on/off	0.9	0.1	−2.7	43%	4%
	nanos3	germ plasm	2.3	9.1	+1.9	50%	44%
	dnd1	germ plasm	2.5	6.2	+1.2	43%	21%
	tbx16	early zygotic	2.5	0.5	−2.2	64%	15%
oblong DS5	*aldh1l1*	near-absent in b	0.7	0.1	−2.3	19%	2%
	chia.6	near-absent in b	0.7	0.0	−2.7	21%	1%
	krt8	EVL	0.0	1.1	+3.0	4%	12%
	krt18	EVL	0.5	3.3	+2.5	7%	21%
	cldne	EVL	0.3	2.8	+3.0	4%	18%
50% epiboly DS4	*hsp90aa1.2*	heat shock	8.4	0.4	−4.0	81%	14%
	hsp70.3	heat shock	1.4	0.1	−3.1	34%	3%
	serpinh1b	heat shock	4.4	0.5	−2.9	57%	14%
	si:dkey-108k21.10	on/off	1.8	0.1	−3.2	49%	5%
bud DS2	*myod1*	muscle (adaxial)	0.7	0.1	−2.2	4%	1%
	myl10	muscle (adaxial)	1.7	0.4	−2.0	12%	8%
	actc1a	muscle (adaxial)	0.7	0.2	−1.6	13%	5%
	sox19b	early neuroectoderm	0.2	0.5	+1.0	10%	21%

References

[1] Maren Büttner, Zhichao Miao, F. Alexander Wolf, Sarah A. Teichmann, and Fabian J. Theis. A test metric for assessing single-cell RNA-seq batch correction. *Nature Methods*, 16(1):43–49, 2019.

[2] Jeffrey A. Farrell, Yiqun Wang, Samantha J. Riesenfeld, Karthik Shekhar, Aviv Regev, and Alexander F. Schier. Single-cell reconstruction of developmental trajectories during zebrafish embryogenesis. *Science*, 360(6392):eaar3131, 2018.

[3] David Gerard. Data-based RNA-seq simulations by binomial thinning. *BMC Bioinformatics*, 21:206, 2020.

[4] Ilya Korsunsky, Nghia Millard, Jean Fan, Kamil Slowikowski, Fan Zhang, Kevin Wei, Yuriy Baglaenko, Michael Brenner, Po-Ru Loh, and Soumya Raychaudhuri. Fast, sensitive and accurate integration of single-cell data with Harmony. *Nature Methods*, 16(12):1289–1296, 2019.

[5] David Lopez-Paz and Maxime Oquab. Revisiting classifier two-sample tests. In *International Conference on Learning Representations (ICLR)*, 2017.

[6] Malte D. Luecken, Maren Büttner, Kridsadakorn Chaichoompu, Anna Danese, Marta Interlandi, Michaela F. Mueller, Daniel C. Strobl, Luke Zappia, Martin Dugas, Maria Colomé-Tatché, and Fabian J. Theis. Benchmarking atlas-level data integration in single-cell genomics. *Nature Methods*, 19(1):41–50, 2022.